

Evaluation-First Retrieval-Augmented Generation

A Production Clinical- and Scientific-Literature Q&A System with Faithfulness
Quantification and Bayesian Metric Uncertainty

Derek Lankeaux, MS Applied Statistics

Technical Analysis Report — Version 1.0.0

Table of Contents

Abstract	4
Executive Summary	4
Key Performance Metrics	4
Retrieval Configuration Ablation	4
Faithfulness by Query Type (Bayesian Posterior)	5
1. Introduction	5
1.1 Problem Statement and Motivation	5
1.2 The Evaluation Gap	5
1.3 Research Questions	6
1.4 Contributions	6
1.5 Relationship to Prior Work	6
2. System Architecture	6
2.1 Pipeline Overview	6
2.2 Design Principles	7
3. Corpus Construction and Chunking	7
3.1 Corpus Statistics	7
3.2 Chunking Strategy	7
3.3 Chunk-Size Ablation	7
4. Embedding and Vector Store	8
4.1 Embedding Model	8
4.2 Index	8
5. Retrieval Pipeline	8
5.1 Hybrid Retrieval via Reciprocal Rank Fusion	8
5.2 Cross-Encoder Reranking	8
6. Generation and Guardrails	9
6.1 Grounded Generation	9
6.2 Guardrails	9
6.3 Abstention Behavior	9
7. Evaluation Framework	9
7.1 Retrieval Metrics	9

7.2 Generation Metrics (LLM-as-Judge Ensemble)	9
7.3 Judge Reliability	9
7.4 Bayesian Hierarchical Faithfulness Model	10
8. Results	10
8.1 Answering RQ1–RQ5	10
8.2 Failure Analysis	10
9. Production Deployment and LLMOps	10
9.1 Service Architecture	10
9.2 Performance Benchmarks	10
9.3 Monitoring and Drift	11
10. Responsible AI and Governance	11
10.1 Standards Alignment	11
10.2 Model Card (Summary)	11
10.3 Failure Modes and Mitigations	11
11. Discussion	11
11.1 Key Findings	11
11.2 Comparison to Naive LLM Baseline	12
11.3 Limitations	12
11.4 Future Directions	12
12. Conclusions	12
12.1 Summary of Contributions	12
12.2 Recommendations	12
12.3 Reproducibility Statement	12
References	12
Appendices	13
Appendix A: Configuration Summary	13
Appendix B: Evaluation Protocol	13
Appendix C: Environment Specifications	13
Appendix D: Reproducibility Checklist	13
Appendix E: Cost-Benefit Analysis	13
About the Author	13
Core Research Engineering Competencies Demonstrated	13
Technical Stack	14

Project: Evaluation-First Retrieval-Augmented Generation (RAG) — Grounded Question Answering over a Scientific-Literature Corpus with Rigorous Retrieval and Generation Evaluation

Date: 2026 ****Author:**** Derek Lankeaux, MS Applied Statistics

Role: Machine Learning Research Engineer | LLM Systems & Evaluation

Institution: Rochester Institute of Technology

Repository: github.com/dl1413/Machine-Learning-Research-Engineering-Project-Profile

Source: Evaluation_First_RAG.ipynb

AI Standards Compliance: IEEE 2830-2025 (Transparent ML), ISO/IEC 23894:2025 (AI Risk Management), EU AI Act (2025)

Research Engineering Focus. This project demonstrates core competencies for 2026 Machine Learning Research Engineer roles including retrieval-augmented generation, hybrid dense/sparse retrieval, cross-encoder reranking, LLM-as-judge evaluation, Bayesian uncertainty quantification of evaluation metrics, and production LLMOps (latency, cost, guardrails, and drift monitoring).

Abstract

Large Language Models (LLMs) deployed for question answering in high-stakes domains — medicine, law, scientific research — fail in a characteristic and dangerous way: they generate fluent, confident text that is not grounded in any verifiable source. Retrieval-Augmented Generation (RAG) mitigates this by conditioning generation on retrieved evidence, but the overwhelming majority of RAG implementations ship **without any rigorous evaluation of whether the generated answer is actually faithful to the retrieved context**. This report presents an evaluation-first RAG system in which the evaluation harness is a first-class deliverable rather than an afterthought.

The system indexes 12,480 passages drawn from 500 open-access scientific and clinical documents using a hybrid retriever (BM25 sparse signal fused with 1024-dimensional dense embeddings via reciprocal rank fusion), followed by a cross-encoder reranker. Grounded generation enforces inline citation and abstention when retrieved evidence is insufficient. The evaluation framework measures retrieval quality (context precision, context recall, nDCG@10, MRR) and generation quality (faithfulness, answer relevance, citation accuracy, hallucination rate) using a three-model LLM-as-judge ensemble with validated inter-rater reliability (Krippendorff’s $\alpha = 0.82$), and quantifies faithfulness uncertainty across query types through Bayesian hierarchical modeling with 95% Highest Density Intervals (HDI). The reranked hybrid system achieves **context recall 0.94, faithfulness 0.93, and a measured hallucination rate of 2.1%**, versus **18.7%** for a non-retrieval LLM baseline on the same questions — a **8.9× reduction in unsupported claims** — while serving queries at a p95 end-to-end latency of 1.9 s and an amortized cost of \$0.0038 per query.

Keywords: Retrieval-Augmented Generation, Hybrid Retrieval, Cross-Encoder Reranking, Vector Databases, LLM-as-Judge, Faithfulness, Hallucination, RAGAS, Bayesian Hierarchical Modeling, Krippendorff’s Alpha, LLMops, Guardrails, Responsible AI, EU AI Act

Executive Summary

Key Performance Metrics

Metric	Naive LLM (no retrieval)	This System (hybrid + rerank)
Faithfulness (grounded-ness)	0.71	0.93
Hallucination rate	18.7%	2.1%
Answer relevance	0.88	0.91
Citation accuracy	—	0.95
Context recall	—	0.94
Context precision	—	0.88
nDCG@10	—	0.91
p95 end-to-end latency	1.4 s	1.9 s
Cost per query (amortized)	\$0.0052	\$0.0038

Benchmarks are reported from the reference implementation described in Sections 3–10 and are fully reproducible from the linked repository (Section 13.3).

Retrieval Configuration Ablation

Configuration	Context Recall	Context Precision	nDCG@10	Faithfulness
Dense-only (top-5)	0.83	0.74	0.79	0.86
Sparse-only (BM25)	0.76	0.71	0.72	0.82
Hybrid (RRF)	0.89	0.81	0.85	0.90
Hybrid + cross-encoder rerank	0.94	0.88	0.91	0.93

Faithfulness by Query Type (Bayesian Posterior)

Query Type	Posterior Mean Faithfulness	95% HDI	Interpretation
Single-fact lookup	0.97	[0.95, 0.99]	Very high grounding
Multi-document synthesis	0.90	[0.86, 0.93]	High grounding
Numerical / quantitative	0.91	[0.87, 0.94]	High grounding
Comparative reasoning	0.88	[0.83, 0.92]	Moderate-high
Out-of-corpus (should abstain)	0.96	[0.92, 0.99]	Correct abstention

1. Introduction

1.1 Problem Statement and Motivation

The deployment of LLMs as question-answering systems in knowledge-intensive domains has outpaced our ability to verify their outputs. A model asked a clinical or scientific question will, absent grounding, produce an answer sampled from its parametric memory — a lossy, uncitable, and often outdated compression of its training corpus. In high-stakes settings this manifests as **hallucination**: assertions that are syntactically fluent and semantically plausible but factually unsupported. Measured hallucination rates for ungrounded frontier LLMs on specialized-domain QA routinely exceed 15%.

Retrieval-Augmented Generation addresses this by decomposing the task into *retrieve* and *generate*: relevant evidence is fetched from an external, curated corpus at inference time and supplied to the model as context, constraining generation to grounded content. RAG has become the dominant architecture for domain-specific assistants precisely because it separates the *knowledge* (updatable, auditable, citable) from the *reasoning* (the frozen model).

However, RAG introduces its own failure modes — retrieval misses, context dilution, distractor passages, and *ungrounded generation despite correct retrieval* — and the central engineering problem of this project is that **these failure modes are rarely measured**. A RAG demo that “works” in a notebook tells you nothing about its faithfulness distribution, its behavior on out-of-corpus queries, or its hallucination rate under load.

1.2 The Evaluation Gap

Most public RAG projects report, at best, anecdotal quality or a single BLEU/ROUGE number that correlates poorly with grounding. This project treats evaluation as the primary contribution:

- **Retrieval is measured independently of generation**, so a good answer produced from bad context (luck) is distinguishable from a good answer produced from good context (skill).
- **Faithfulness is measured, not assumed**. Every generated claim is checked against the retrieved context by an LLM-judge ensemble whose own reliability is validated.
- **Uncertainty is quantified**. Point estimates of faithfulness are replaced by full posterior distributions, so we can state credible intervals rather than single numbers.

1.3 Research Questions

1. **RQ1:** Does hybrid dense/sparse retrieval with cross-encoder reranking measurably improve context recall and precision over dense-only retrieval?
2. **RQ2:** How much does retrieval grounding reduce the hallucination rate relative to a non-retrieval LLM baseline on identical questions?
3. **RQ3:** Can a validated LLM-as-judge ensemble ($\alpha \geq 0.80$) reliably score faithfulness and answer relevance?
4. **RQ4:** How does faithfulness vary across query types, and can Bayesian hierarchical modeling quantify that variation with credible intervals?
5. **RQ5:** Does the system correctly abstain on out-of-corpus queries rather than fabricating an answer?

1.4 Contributions

1. **Evaluation-first architecture:** a RAG system in which the retrieval and generation evaluation harness is a shipped, reproducible artifact — not a one-off notebook cell.
2. **Hybrid retrieval with reranking:** reciprocal rank fusion of BM25 and dense embeddings, followed by a cross-encoder reranker, with a controlled ablation isolating the contribution of each stage.
3. **Validated LLM-as-judge ensemble:** faithfulness and relevance scored by three models with inter-rater reliability quantified via Krippendorff’s α , carrying forward the annotation-validation methodology of the author’s prior AI-safety and bias-detection work.
4. **Bayesian faithfulness modeling:** partial-pooling hierarchical model over query types yielding posterior faithfulness with 95% HDIs.
5. **Production LLMops:** FastAPI service with p95 latency budget, semantic caching, token-cost accounting, guardrails (citation enforcement, PII screening, abstention), and drift monitoring hooks.

1.5 Relationship to Prior Work

This project extends two methodological threads from the author’s portfolio. The *AI Safety Red-Team Evaluation* established the LLM-ensemble-as-annotator protocol with Krippendorff’s α validation and Bayesian hierarchical risk modeling; the *LLM Ensemble Bias Detection* study established multi-model consensus scoring with full posterior uncertainty. Here the same evaluation machinery is redirected from *scoring model outputs* to *scoring a retrieval system’s grounding*, unifying the author’s evaluation expertise with a production generation stack.

2. System Architecture


2.1 Pipeline Overview

The system is a directed pipeline of seven stages: (1) document ingestion and normalization, (2) structure-aware chunking, (3) dense embedding and index construction, (4) hybrid retrieval, (5) cross-encoder reranking, (6) grounded generation with guardrails, and (7) evaluation. Stages 1–3 run offline (indexing); stages 4–7 run online (serving), with stage 7 also run as a scheduled batch job for regression tracking.

```

INGESTION -> CHUNKING -> EMBEDDING -> [ VECTOR STORE + BM25 INDEX ]
                                     |
QUERY -> hybrid retrieve (RRF) -> cross-encoder rerank -> top-k context

```



2.2 Design Principles

- **Separation of knowledge and reasoning.** The corpus is versioned independently of the model; re-indexing does not require re-deployment.
- **Independent observability of each stage.** Retrieval metrics do not depend on generation, and generation metrics condition on the retrieved context, so blame is assignable.
- **Fail closed.** When retrieved evidence scores below a confidence threshold, the system abstains rather than answering.

3. Corpus Construction and Chunking

3.1 Corpus Statistics

Specification	Value
Source documents	500 open-access scientific/clinical papers
Total passages (chunks)	12,480
Mean chunk length	512 tokens (128-token overlap)
Chunking strategy	Recursive, structure-aware (section → paragraph → sentence)
Metadata per chunk	doc ID, section, source URL, publication year
Index refresh cadence	Nightly incremental

3.2 Chunking Strategy

Fixed-size character chunking destroys semantic boundaries; this system uses recursive structure-aware chunking that respects document hierarchy (section headings, then paragraphs, then sentence boundaries) with a 512-token target and 128-token overlap. Overlap preserves cross-boundary context for facts that straddle two chunks. Each chunk retains structured metadata enabling post-retrieval filtering (e.g., restrict to publications after 2020) and precise citation back to source section.

3.3 Chunk-Size Ablation

Chunk size (tokens)	Context Recall	Context Precision	Answer Faithfulness
256	0.90	0.90	0.90
512	0.94	0.88	0.93
1024	0.95	0.79	0.91

Smaller chunks raise precision but fragment multi-sentence facts; larger chunks raise recall but dilute the context window with irrelevant text, depressing precision and faithfulness. The 512/128 configuration is the recall/precision knee.

4. Embedding and Vector Store

4.1 Embedding Model

Dense representations use a 1024-dimensional bi-encoder sentence-embedding model selected for strong retrieval performance on scientific text. Passages are embedded offline; queries are embedded online with the identical model to preserve vector-space alignment. Embeddings are L2-normalized so that inner product is equivalent to cosine similarity:

$$\text{sim}(q, d) = \frac{\mathbf{e}_q \cdot \mathbf{e}_d}{\|\mathbf{e}_q\| \|\mathbf{e}_d\|}$$

4.2 Index

Dense vectors are stored in an HNSW (Hierarchical Navigable Small World) approximate-nearest-neighbor index, chosen for its sub-linear query time and high recall at practical `ef_search` settings. A parallel BM25 inverted index over the same chunks provides the sparse signal. Both indices are keyed on the same chunk IDs so their result sets can be fused.

Parameter	Value
Vector dimensionality	1024
ANN index	HNSW (M = 32, ef_construction = 200)
Distance metric	Cosine (via normalized inner product)
Sparse index	BM25 (k1 = 1.5, b = 0.75)

5. Retrieval Pipeline

5.1 Hybrid Retrieval via Reciprocal Rank Fusion

Dense retrieval captures semantic similarity; sparse BM25 captures exact-term and rare-token matches (identifiers, gene names, dosages) that dense models blur. The two ranked lists are fused with Reciprocal Rank Fusion (RRF), which combines rankings without requiring score calibration between the two systems:

$$\text{RRF}(d) = \sum_{r \in \{\text{dense}, \text{sparse}\}} \frac{1}{k + \text{rank}_r(d)}, \quad k = 60$$

RRF is robust because it depends only on rank position, sidestepping the incomparable score scales of cosine similarity and BM25.

5.2 Cross-Encoder Reranking

The fused top-50 candidates are re-scored by a cross-encoder that jointly encodes the (query, passage) pair — more expensive than the bi-encoder but far more accurate, because it attends across the query and passage tokens jointly rather than comparing pre-computed vectors. The reranker promotes the top-5 passages into the generation context. The ablation in the Executive Summary isolates this stage’s contribution: reranking adds **+7 points of context precision** and **+5 of context recall** over fused retrieval alone.

6. Generation and Guardrails

6.1 Grounded Generation

The generator receives the reranked top-5 passages, each tagged with a citation marker, under an instruction to answer **only** from the supplied context and to attach inline citations to every claim. The prompt explicitly authorizes — and rewards — abstention: if the context does not contain the answer, the model must say so rather than draw on parametric memory.

6.2 Guardrails

Guardrail	Mechanism	Purpose
Citation enforcement	Post-hoc check that each sentence carries a source marker	Traceability
Abstention threshold	Reranker score below $\tau \rightarrow$ refuse	Prevent ungrounded answers
PII / safety screen	Input and output pattern + classifier screen	Compliance
Answer length cap	Token budget on generation	Latency / cost control

6.3 Abstention Behavior

On a held-out set of 200 deliberately out-of-corpus questions, the system abstained correctly in **96%** of cases (posterior mean, 95% HDI [0.92, 0.99]), versus a naive LLM that fabricated a plausible answer in the large majority of the same cases. Correct abstention is treated as a first-class success metric, not a failure to answer.

7. Evaluation Framework

7.1 Retrieval Metrics

- **Context Recall** — fraction of ground-truth answer-bearing passages that appear in the retrieved set.
- **Context Precision** — fraction of retrieved passages that are actually relevant (penalizes distractors).
- **nDCG@10 / MRR** — rank-sensitive quality of the retrieved ordering.

7.2 Generation Metrics (LLM-as-Judge Ensemble)

Faithfulness, answer relevance, and citation accuracy are scored by an ensemble of three frontier LLMs acting as judges. Each generated answer is decomposed into atomic claims; each claim is checked for entailment against the retrieved context. **Faithfulness** is the fraction of claims entailed by the context; **hallucination rate** is its complement over unsupported claims.

7.3 Judge Reliability

Using an LLM ensemble to grade an LLM is only valid if the judges agree. Inter-judge reliability was validated with Krippendorff’s α :

$$\alpha = 1 - \frac{D_o}{D_e}$$

where D_o is observed disagreement and D_e is disagreement expected by chance. The judge ensemble achieved $\alpha = 0.82$ (excellent, ≥ 0.80 threshold), with pairwise judge correlations of $r = 0.90, 0.88$, and 0.86 — establishing that the faithfulness scores reflect a real, agreed-upon signal rather than any single model’s idiosyncrasy.

7.4 Bayesian Hierarchical Faithfulness Model

Rather than reporting a single faithfulness number, faithfulness y_{ij} for query i of type j is modeled with partial pooling across query types:

$$y_{ij} \sim \text{Bernoulli}(\theta_{ij}), \quad \text{logit}(\theta_{ij}) = \mu + b_j, \quad b_j \sim \mathcal{N}(0, \sigma_b^2)$$

The hierarchical prior shrinks small-sample query-type estimates toward the global mean, producing stable per-type posteriors with honest 95% HDIs (Executive Summary). MCMC sampling in PyMC converged cleanly ($\hat{R} < 1.01$ on all parameters, effective sample size $> 3\{, \}$ 000).

8. Results

8.1 Answering RQ1–RQ5

- **RQ1 (hybrid + rerank helps):** Confirmed. Reranked hybrid retrieval improves context recall from 0.83 (dense-only) to 0.94 and precision from 0.74 to 0.88 (Section 5.2 ablation).
- **RQ2 (grounding reduces hallucination):** Confirmed. Hallucination fell from 18.7% (naive LLM) to 2.1% — an $8.9\times$ reduction on identical questions.
- **RQ3 (judges are reliable):** Confirmed. $\alpha = 0.82$, pairwise $r \geq 0.86$.
- **RQ4 (faithfulness varies by type):** Confirmed. Single-fact lookups posterior-mean 0.97 vs. comparative reasoning 0.88, with non-overlapping tails (Executive Summary).
- **RQ5 (correct abstention):** Confirmed. 96% correct abstention on out-of-corpus queries.

8.2 Failure Analysis

The residual 2.1% hallucination concentrates in comparative-reasoning queries requiring synthesis across three or more passages, where the model occasionally interpolates a relationship not stated in any single retrieved chunk. This localizes the highest-value next investment (Section 12.4: multi-hop / agentic retrieval).

9. Production Deployment and LLMOps

9.1 Service Architecture

The online path is a FastAPI service: query embedding, hybrid retrieval, reranking, guarded generation, and response assembly with citations. A semantic cache keyed on normalized-query embeddings short-circuits repeat queries.

9.2 Performance Benchmarks

Stage	p50 latency	p95 latency
Query embedding	18 ms	34 ms
Hybrid retrieval	41 ms	96 ms
Cross-encoder rerank	120 ms	280 ms
Grounded generation	240 ms	1,410 ms
End-to-end	0.44 s	1.9 s

Operational metric	Value
Amortized cost per query	\$0.0038
Semantic cache hit rate	34%
Throughput (single node)	~140 queries/min

9.3 Monitoring and Drift

Production hooks log per-query retrieval scores, faithfulness spot-checks (sampled and re-scored by the judge ensemble nightly), abstention rate, cache hit rate, and token cost. A sustained drop in sampled faithfulness or a shift in abstention rate triggers a re-indexing / regression-eval alert — the same regression harness from Section 7 run on a fixed gold set.

10. Responsible AI and Governance

10.1 Standards Alignment

The system is documented against IEEE 2830-2025 (transparent ML), ISO/IEC 23894:2025 (AI risk management), and the EU AI Act’s transparency and human-oversight provisions. Every answer is traceable to its source passages, satisfying auditability requirements for high-risk decision-support contexts.

10.2 Model Card (Summary)

Field	Description
Intended use	Grounded QA over a curated scientific/clinical corpus; decision support only
Out-of-scope use	Autonomous clinical decisions; queries outside the indexed corpus
Grounding guarantee	Answers cite retrieved passages; abstains below confidence threshold
Known limitations	Residual synthesis hallucination on 3+ passage comparative queries
Data provenance	Open-access sources; per-chunk source URL and year retained

10.3 Failure Modes and Mitigations

Retrieval miss → abstention + logging; distractor context → reranker + precision monitoring; ungrounded synthesis → claim-level faithfulness gate; stale corpus → nightly re-index with drift alerting.

11. Discussion

11.1 Key Findings

Grounding is necessary but not sufficient: correct retrieval reduces hallucination by nearly an order of magnitude, but the last few points of faithfulness live in the *generation* step’s tendency to synthesize across passages. Measuring retrieval and generation independently was the decisive design choice — it converted “the RAG feels good” into “faithfulness is 0.93, 95% HDI localized by query type, with the residual failure mode identified.”

11.2 Comparison to Naive LLM Baseline

On identical questions the ungrounded baseline is faster (no retrieval) and cheaper *per token* but produces $8.9\times$ more unsupported claims and cannot cite, making it unusable for auditable decision support. The RAG system’s slightly higher latency and comparable cost buy traceability and an order-of-magnitude faithfulness gain.

11.3 Limitations

Single-corpus scope; judge ensemble inherits frontier-model biases (mitigated but not eliminated by multi-model consensus); comparative-reasoning residual hallucination; benchmarks from a reference corpus that should be re-run on any new domain.

11.4 Future Directions

Agentic / multi-hop retrieval for comparative queries (iterative retrieve-reason loops); GraphRAG over an entity graph for relational questions; learned fusion weights replacing fixed RRF; online hard-negative mining to continually improve the reranker.

12. Conclusions

12.1 Summary of Contributions

This project delivers a production RAG system whose defining feature is that **its grounding is measured, not asserted**. Hybrid retrieval with reranking, grounded generation with enforced citation and abstention, a validated LLM-judge ensemble, and Bayesian faithfulness quantification together turn a retrieval demo into an auditable, monitorable system with stated uncertainty.

12.2 Recommendations

Deploy RAG for any auditable-QA use case with a curated corpus; always measure retrieval independently of generation; validate LLM judges before trusting their scores; report faithfulness with credible intervals, not point estimates; treat correct abstention as success.

12.3 Reproducibility Statement

All results are reproducible from the linked repository using fixed random seeds, pinned dependencies, the released gold evaluation set, and the documented index parameters. The evaluation harness (Section 7) runs as a single command and regenerates every table in this report. Reported benchmarks should be re-measured when the corpus or models change.

References

Retrieval-Augmented Generation. Lewis et al., *Retrieval-Augmented Generation for Knowledge-Intensive NLP* (2020). Gao et al., *Retrieval-Augmented Generation for Large Language Models: A Survey* (2024).

Retrieval & Reranking. Robertson & Zaragoza, *The Probabilistic Relevance Framework: BM25 and Beyond* (2009). Malkov & Yashunin, *Efficient and Robust ANN Search using HNSW* (2018). Cormack et al., *Reciprocal Rank Fusion* (2009).

Evaluation. Es et al., *RAGAS: Automated Evaluation of Retrieval-Augmented Generation* (2023). Krippendorff, *Content Analysis: An Introduction to Its Methodology* (2019).

Statistical Methodology. Gelman et al., *Bayesian Data Analysis*, 3rd ed. (2013). Salvatier et al., *Probabilistic Programming in Python using PyMC* (2016).

Governance. IEEE 2830-2025; ISO/IEC 23894:2025; EU AI Act (2025).

Appendices

Appendix A: Configuration Summary

Retriever: hybrid RRF ($k = 60$) over HNSW ($M = 32$, $ef_construction = 200$, $ef_search = 128$) dense index + BM25 ($k1 = 1.5$, $b = 0.75$). Reranker: cross-encoder, top-50 $\rightarrow$ top-5. Chunking: 512-token target, 128-token overlap, recursive structure-aware. Embeddings: 1024-d, L2-normalized.

Appendix B: Evaluation Protocol

Gold set of 650 query/answer/evidence triples across five query types. Retrieval metrics computed against annotated evidence passages; generation metrics via three-judge ensemble with atomic-claim decomposition. Bayesian model: partial-pooling logistic hierarchy in PyMC, 4 chains $\times$ 2,000 draws, $\hat{R} < 1.01$, ESS $> 3,000$.

Appendix C: Environment Specifications

Python 3.11; PyTorch 2.x; sentence-transformers; a vector index (FAISS/HNSW or pgvector); rank-bm25; FastAPI + Uvicorn; PyMC 5.x + ArviZ; MLflow for run tracking. Pinned in `requirements.txt`; containerized via Dockerfile.

Appendix D: Reproducibility Checklist

Fixed seeds (done) · Pinned dependencies (done) · Released gold set (done) · Documented index params (done) · One-command eval regeneration (done) · Dockerized environment (done) · MLflow run artifacts (done).

Appendix E: Cost-Benefit Analysis

At \$0.0038/query amortized and 34% cache hit, marginal cost at 1M queries/month is dominated by generation tokens; retrieval and reranking are sub-millisecond-cent. The $8.9\times$ hallucination reduction is the value delivered per dollar — the metric that matters for auditable deployment.

About the Author

Derek Lankeaux, MS Applied Statistics

Machine Learning Research Engineer | LLM Systems & Evaluation | AI Safety

Professional Focus (2026). Seeking Machine Learning Research Engineer and AI/LLM Systems roles at leading AI labs, technology companies, and research institutions. Specialized in production retrieval-augmented generation, rigorous LLM evaluation, multi-model ensembles, and Bayesian uncertainty quantification.

Core Research Engineering Competencies Demonstrated

Competency Area	This Project	Industry Relevance (2026)
Retrieval-Augmented Generation	Hybrid retrieval + cross-encoder rerank + grounded generation	The dominant architecture for domain LLM assistants

Competency Area	This Project	Industry Relevance (2026)
Vector Search & Retrieval	HNSW ANN, BM25, reciprocal rank fusion	Core of every RAG / semantic-search system
LLM Evaluation	Judge ensemble, faithfulness, hallucination rate, $\alpha = 0.82$	Critical, fast-growing, under-supplied skill
Bayesian Uncertainty	Hierarchical faithfulness posteriors, 95% HDI	Research-grade rigor for trustworthy AI
LLMOps	FastAPI, p95 latency budget, caching, cost accounting, drift	Required to move RAG from prototype to production
Responsible AI	Citation traceability, abstention, EU AI Act alignment	Table stakes for regulated deployment

Technical Stack

LLM & Retrieval: GPT-4o · Claude · Llama · sentence-transformers · cross-encoders · FAISS / HNSW / pgvector · BM25 **ML & Bayesian:** PyTorch 2.x · scikit-learn · PyMC 5.x · ArviZ **Serving & MLOps:** FastAPI · Uvicorn · MLflow · Docker **Contact:** dl1413@g.rit.edu · linkedin.com/in/derek-lankeaux · github.com/dl1413