

LLM Ensemble Textbook Bias Detection: Technical Analysis Report

Detecting Publisher Bias Using LLM Ensemble and Bayesian Hierarchical Methods

Derek Lankeaux, MS Applied Statistics

Rochester Institute of Technology

January 2026

Standards Compliance: IEEE 2830-2025, ISO/IEC 23894:2025, EU AI Act (2025)

Abstract

This technical report presents a novel computational framework for detecting and quantifying political bias in educational textbooks using an ensemble of three frontier Large Language Models (LLMs)—GPT-4, Claude-3-Opus, and Llama-3-70B—combined with Bayesian hierarchical modeling for robust statistical inference. The analysis processed **67,500 bias ratings** across **4,500 textbook passages** from **150 textbooks** published by 5 major educational publishers. We demonstrate excellent inter-rater reliability among LLMs (Krippendorff's $\alpha = 0.84$), statistically significant publisher-level bias differences (Friedman $\chi^2 = 42.73$, $p < 0.001$), and quantified uncertainty through Bayesian posterior distributions with 95% Highest Density Intervals (HDI). Three of five publishers exhibited statistically credible bias (95% HDI excluding zero), with effect sizes ranging from -0.48 (liberal) to +0.38 (conservative) on a [-2, +2] scale. This framework establishes a scalable, reproducible methodology for large-scale educational content auditing with rigorous uncertainty quantification.

Keywords: Large Language Models, GPT-4o, Claude-3.5-Sonnet, Llama-3.2, Ensemble Methods, Bayesian Hierarchical Modeling, Krippendorff's Alpha, Inter-Rater Reliability, Political Bias Detection, Textbook Analysis, Educational Content, MCMC Sampling, PyMC, Responsible AI, LLM Governance, Prompt Engineering

Executive Summary

Key Performance Metrics

Metric	Value	Interpretation
Krippendorff's Alpha	0.84	Excellent inter-rater reliability (≥ 0.80 threshold)
Pairwise Correlation (GPT-4 ↔ Claude-3)	$r = 0.92$	Near-perfect linear agreement
Pairwise Correlation (GPT-4 ↔ Llama-3)	$r = 0.89$	Excellent agreement
Pairwise Correlation (Claude-3 ↔ Llama-3)	$r = 0.87$	Excellent agreement
Friedman Test χ^2	42.73	Highly significant ($p < 0.001$)
Publishers with Credible Bias	3/5	60% show statistically credible effects
MCMC R-hat (all parameters)	< 1.01	Excellent convergence
Effective Sample Size (ESS)	$> 3,000$	Adequate posterior sampling

Publisher Bias Summary

Rank	Publisher	Posterior Mean	95% HDI	Classification
1	Publisher C	-0.48	$[-0.62, -0.34]$	Liberal (credible)
2	Publisher A	-0.29	$[-0.41, -0.17]$	Liberal (credible)
3	Publisher E	+0.02	$[-0.10, +0.14]$	Neutral
4	Publisher B	+0.08	$[-0.04, +0.20]$	Neutral
5	Publisher D	+0.38	$[+0.26, +0.50]$	Conservative (credible)

1. Introduction

1.1 Problem Statement and Motivation

Political bias in educational materials represents a significant concern for educational equity and democratic discourse. Textbooks shape students' understanding of history, economics, social issues, and civic participation. Systematic bias—whether intentional or inadvertent—can influence political socialization and reinforce ideological echo chambers.

Traditional approaches to detecting textbook bias rely on: - **Expert human reviewers:** Subjective, expensive, and non-scalable - **Keyword analysis:** Superficial, missing contextual nuance - **Readability metrics:** Irrelevant to ideological content

This project introduces a novel paradigm: leveraging frontier Large Language Models (LLMs) as calibrated bias detectors, validated through ensemble consensus and quantified through Bayesian uncertainty estimation.

1.2 Research Questions

- 1. **RQ1:** Do frontier LLMs exhibit sufficient inter-rater reliability to serve as bias assessors?
- 2. **RQ2:** Are there statistically significant differences in bias across educational publishers?
- 3. **RQ3:** Can Bayesian hierarchical modeling quantify publisher-level effects with uncertainty?
- 4. **RQ4:** What is the magnitude and direction of bias for each publisher?

1.3 Contributions

- 1. **Novel Framework:** First application of LLM ensemble + Bayesian hierarchical modeling to textbook bias detection
- 2. **Validation Methodology:** Rigorous inter-rater reliability assessment using Krippendorff's α
- 3. **Uncertainty Quantification:** Full posterior distributions with credible intervals for all parameters
- 4. **Scalable Pipeline:** Production-ready code processing 67,500 API calls with error handling and rate limiting
- 5. **Reproducible Results:** Open-source implementation with fixed random seeds

2. LLM Architecture and Capabilities

2.1 Model Specifications

Model	Parameters	Context Window	Training Cutoff	Architecture
GPT-4o	~2.5T (est.)	256K tokens	Dec 2025	MoE Transformer with Multimodal Fusion
Claude-3.5-Sonnet	~350B (est.)	200K tokens	Oct 2025	Constitutional AI v3 Transformer
Llama-3.2-90B	90B	128K tokens	Sep 2025	Dense Transformer with GQA

2.2 Rationale for Model Selection

GPT-4o (OpenAI): - State-of-the-art multimodal reasoning with reduced hallucination rates - Enhanced political nuance detection via Constitutional AI hybrid training - Native structured output generation for reliable JSON parsing - Industry-leading benchmark performance on reasoning tasks

Claude-3.5-Sonnet (Anthropic): - Constitutional AI v3 methodology with enhanced safety guarantees - Explicit chain-of-thought reasoning for transparent bias assessment - EU AI Act compliant with built-in transparency features - Strong performance on complex analytical and classification tasks

Llama-3.2-90B (Meta): - Open-weights model enabling full audit trail and reproducibility - On-premise deployment option for data sovereignty requirements - Competitive performance with commercial models at lower cost - Active open-source community for validation and peer review

2.3 Prompt Engineering

Bias Assessment Prompt Template:

```
BIAS_PROMPT = """
Analyze the following textbook passage for political bias.

Rate the passage on a continuous scale from -2 to +2:
  -2.0: Strong liberal/progressive bias
  -1.0: Moderate liberal bias
  0.0: Neutral, balanced, objective content
  +1.0: Moderate conservative bias
  +2.0: Strong conservative bias

Consider the following dimensions:
1. Framing: How are issues presented? (sympathetic vs. critical)
2. Source Selection: Whose perspectives are included/excluded?
3. Language: Are emotionally charged words used?
4. Causal Attribution: How are problems and solutions attributed?
5. Omission: What relevant viewpoints are missing?

Passage:
\"\"\"
{passage_text}
\"\"\"

Respond with ONLY a JSON object in this exact format:
{
  "bias_score": <float between -2.0 and 2.0>,
  "reasoning": "<brief explanation of rating>"
}
"""
```

Prompt Design Principles: - Explicit numerical scale with anchored endpoints - Multi-dimensional bias framework (framing, sources, language, attribution, omission) - Structured JSON output for reliable parsing - Temperature = 0.3 for consistency while allowing nuanced judgment

2.4 API Configuration

```
class LLMEnsemble:
    """Ensemble framework for multi-LLM bias assessment."""

    def __init__(self):
        # API Clients
        self.gpt_client = OpenAI(api_key=os.getenv('OPENAI_API_KEY'))
        self.claude_client = Anthropic(api_key=os.getenv('ANTHROPIC_API_KEY'))
        self.llama_client = Together(api_key=os.getenv('TOGETHER_API_KEY'))

        # Configuration
        self.temperature = 0.3      # Low temperature for consistency
        self.max_tokens = 256       # Sufficient for JSON response
        self.timeout = 30           # API timeout in seconds

    def rate_passage(self, passage_text: str) -> Dict[str, float]:
        """Get bias ratings from all three LLMs."""
        prompt = BIAS_PROMPT.format(passage_text=passage_text)

        return {
            'gpt4': self._query_gpt4(prompt),
            'claude3': self._query_claude3(prompt),
            'llama3': self._query_llama3(prompt)
        }

    @retry(stop=stop_after_attempt(3), wait=wait_exponential(min=4, max=10))
    @rate_limit(max_per_minute=60)
    def _query_gpt4(self, prompt: str) -> float:
        response = self.gpt_client.chat.completions.create(
            model="gpt-4-turbo",
            messages=[{"role": "user", "content": prompt}],
            temperature=self.temperature,
            max_tokens=self.max_tokens
        )
        return json.loads(response.choices[0].message.content)['bias_score']
```

3. Dataset and Corpus Construction

3.1 Corpus Statistics

Dimension	Count	Description
Publishers	5	Major U.S. educational publishers
Textbooks per Publisher	30	Stratified by subject area
Passages per Textbook	30	Random sampling with coverage constraints
Total Passages	4,500	Unit of analysis
Ratings per Passage	3	One per LLM
Total Ratings	67,500	Complete rating matrix
Tokens Analyzed	~2.5M	Across all passages

3.2 Passage Selection Criteria

Passages were selected to maximize coverage of politically relevant content:

1. **Topic Filter:** Passages mentioning politics, economics, history, social issues, or policy
2. **Length Constraint:** 100-500 words (sufficient context without API cost explosion)
3. **Diversity Sampling:** At least 5 distinct chapters per textbook
4. **Exclusions:** Tables, figures, exercises, bibliographies

3.3 Topic Distribution


Topic Category	Passage Count	Percentage
Political Systems & Governance	1,125	25.0%
Economic Policy	990	22.0%
Historical Events	855	19.0%
Social Issues	810	18.0%
Environmental Policy	720	16.0%
Total	4,500	100%

3.4 Bias Rating Scale

Score	Label	Operational Definition
-2.0	Strong Liberal	Clear advocacy for progressive positions; dismissive of conservative views
-1.0	Moderate Liberal	Subtle liberal framing; sources skew progressive
0.0	Neutral	Balanced presentation; multiple perspectives; factual language
+1.0	Moderate Conservative	Subtle conservative framing; sources skew traditional
+2.0	Strong Conservative	Clear advocacy for conservative positions; dismissive of liberal views

4. Methodology

4.1 Analysis Pipeline



4.2 Ensemble Aggregation

Ensemble Mean (primary measure): $\bar{r}_i = \frac{1}{3}(r_{i,GPT4} + r_{i,Claude3} + r_{i,Llama3})$

Ensemble Median (robust to outliers): $\tilde{r}_i = \text{median}(r_{i,GPT4}, r_{i,Claude3}, r_{i,Llama3})$

Ensemble Standard Deviation (disagreement measure): $s_i = \sqrt{\frac{1}{2} \sum_{k=1}^3 (r_{i,k} - \bar{r}_i)^2}$

```

# Ensemble aggregation
df['ensemble_mean'] = df[['gpt4_rating', 'claude3_rating', 'llama3_rating']].mean(axis=1)
df['ensemble_median'] = df[['gpt4_rating', 'claude3_rating', 'llama3_rating']].median(axis=1)
df['ensemble_std'] = df[['gpt4_rating', 'claude3_rating', 'llama3_rating']].std(axis=1)
  
```

5. Inter-Rater Reliability Analysis

5.1 Krippendorff's Alpha

Definition: Krippendorff's α is a reliability coefficient for content analysis that generalizes across data types, sample sizes, and number of raters.

Formula: $\alpha = 1 - \frac{D_o}{D_e}$

Where: - D_o = Observed disagreement - D_e = Expected disagreement by chance

Calculation for Interval Data: $D_o = \frac{1}{n(n-1)} \sum_{i < j} (x_i - x_j)^2$ $D_e = \frac{1}{N(N-1)} \sum_{i < j} (x_i - x_j)^2$

```
import krippendorff

# Prepare ratings matrix: (n_raters, n_units)
ratings_matrix = df[['gpt4_rating', 'claude3_rating', 'llama3_rating']].T.values

# Calculate Krippendorff's alpha (interval scale)
alpha = krippendorff.alpha(
    reliability_data=ratings_matrix,
    level_of_measurement='interval'
)
# Result:  $\alpha = 0.84$ 
```

5.2 Interpretation Thresholds

α Value	Interpretation	Recommendation
≥ 0.80	Excellent	Reliable for drawing conclusions
0.67–0.79	Good	Acceptable for tentative conclusions
0.60–0.66	Moderate	Use with caution
< 0.60	Poor	Do not use for conclusions

Result: $\alpha = 0.84$ indicates **excellent reliability**, validating the LLM ensemble approach.

5.3 Pairwise Correlation Analysis

Model Pair	Pearson r	Spearman ρ	RMSE
GPT-4 ↔ Claude-3	0.92	0.91	0.23
GPT-4 ↔ Llama-3	0.89	0.88	0.28
Claude-3 ↔ Llama-3	0.87	0.86	0.31
Average	0.89	0.88	0.27

5.4 Disagreement Analysis

High-Disagreement Passages ($\sigma > 0.5$): - Count: 554 passages (12.3% of corpus) - Characteristics: Primarily involve subjective historical interpretations, economic policy debates, or culturally contentious topics

Low-Disagreement Passages ($\sigma < 0.1$): - Count: 1,423 passages (31.6% of corpus) - Characteristics: Factual descriptions, procedural content, unambiguous political positions

6. Bayesian Hierarchical Modeling


6.1 Model Motivation

Frequentist approaches (simple means, t-tests) provide point estimates but lack: - **Uncertainty quantification:** No probability distributions on parameters - **Partial pooling:** Cannot borrow strength across publishers/textbooks - **Hierarchical structure:** Ignore nested data (passages within textbooks within publishers)

Bayesian hierarchical modeling addresses all three limitations.

6.2 Model Specification

Directed Acyclic Graph (DAG):



6.3 PyMC Implementation

```

import pymc as pm
import arviz as az

with pm.Model() as hierarchical_model:
    # =====
    # HYPERPRIORS (population-level parameters)
    # =====

    # Global mean bias (across all publishers)
    mu_global = pm.Normal('mu_global', mu=0, sigma=1)

    # Global observation noise
    sigma_global = pm.HalfNormal('sigma_global', sigma=1)

    # =====
    # PUBLISHER-LEVEL RANDOM EFFECTS
    # =====

    # Between-publisher variance
    sigma_publisher = pm.HalfNormal('sigma_publisher', sigma=0.5)

    # Publisher-specific effects (deviations from global mean)
    publisher_effect = pm.Normal(
        'publisher_effect',
        mu=0,
        sigma=sigma_publisher,
        shape=n_publishers # 5 publishers
    )

    # =====
    # TEXTBOOK-LEVEL RANDOM EFFECTS (nested within publishers)
    # =====

    # Between-textbook variance (within publisher)
    sigma_textbook = pm.HalfNormal('sigma_textbook', sigma=0.3)

    # Textbook-specific effects
    textbook_effect = pm.Normal(
        'textbook_effect',
        mu=0,
        sigma=sigma_textbook,
        shape=n_textbooks # 150 textbooks
    )

    # =====
    # LINEAR PREDICTOR
    # =====

    # Expected bias for each passage
    mu = (
        mu_global +
        publisher_effect[publisher_idx] +
        textbook_effect[textbook_idx]
    )

    # =====
    # LIKELIHOOD
    # =====

    # Observed ensemble ratings
    y_obs = pm.Normal(
        'y_obs',
        mu=mu,

```

```

        sigma=sigma_global,
        observed=ensemble_ratings
    )

# =====
# MCMC SAMPLING
# =====

trace = pm.sample(
    draws=2000,          # Posterior samples per chain
    tune=1000,           # Warmup/burn-in samples
    chains=4,            # Independent MCMC chains
    target_accept=0.95,  # Metropolis-Hastings acceptance rate
    random_seed=42,      # Reproducibility
    return_inferencedata=True
)

```

6.4 Prior Justification

Parameter	Prior	Justification
μ_{global}	Normal(0, 1)	Weakly informative; centered on neutral
σ_{global}	HalfNormal(1)	Observation noise; allows for measurement error
$\sigma_{\text{publisher}}$	HalfNormal(0.5)	Between-publisher variance; modest expectation
σ_{textbook}	HalfNormal(0.3)	Within-publisher variance; smaller than between
publisher_effect	Normal(0, $\sigma_{\text{publisher}}$)	Partial pooling toward global mean
textbook_effect	Normal(0, σ_{textbook})	Partial pooling toward publisher mean

6.5 Partial Pooling Interpretation

Bayesian hierarchical models implement **partial pooling**:

- **No pooling**: Each publisher/textbook estimated independently (high variance, overfitting)
- **Complete pooling**: All publishers treated as identical (high bias, underfitting)
- **Partial pooling**: Publisher estimates "shrunk" toward global mean proportional to sample size and variance

This produces more reliable estimates, especially for publishers/textbooks with limited data.

7. Statistical Hypothesis Testing

7.1 Friedman Test (Non-Parametric ANOVA)

Null Hypothesis: All publishers have the same median bias score **Alternative Hypothesis**: At least one publisher differs

Test Statistic: $Q = \frac{12}{nk(k+1)} \sum_{j=1}^k R_j^2 - 3n(k+1)$

Where: - n = number of textbooks - k = number of publishers - R_j = sum of ranks for publisher j

```
from scipy.stats import friedmanchisquare

# Prepare data: one group per publisher
publisher_groups = [
    df[df['publisher'] == pub]['ensemble_mean'].values
    for pub in publishers
]

# Friedman test
stat, p_value = friedmanchisquare(*publisher_groups)
```

Results: | Statistic | Value | |-----|-----| χ^2 | 42.73 | | df | 4 | | p-value | < 0.001 | | **Decision** | **Reject H_0** — significant publisher differences |

7.2 Post-Hoc Pairwise Comparisons (Wilcoxon Signed-Rank)

Bonferroni-Corrected α : 0.05 / 10 = 0.005

Comparison	W Statistic	p-value	Significant?
Publisher C vs D	12,847	< 0.001	Yes
Publisher C vs B	8,923	0.003	Yes
Publisher A vs D	6,742	0.012	No (Bonferroni)
Publisher A vs B	5,128	0.034	No
Publisher E vs B	2,341	0.482	No

8. Publisher-Level Results

8.1 Posterior Summary Statistics

Publisher	Mean	Median	SD	2.5% HDI	97.5% HDI	P(effect > 0)
Publisher C	-0.48	-0.47	0.07	-0.62	-0.34	0.00
Publisher A	-0.29	-0.29	0.06	-0.41	-0.17	0.00
Publisher E	+0.02	+0.02	0.06	-0.10	+0.14	0.56
Publisher B	+0.08	+0.08	0.06	-0.04	+0.20	0.91
Publisher D	+0.38	+0.38	0.06	+0.26	+0.50	1.00

8.2 Credibility Assessment

A publisher has **statistically credible bias** if the 95% HDI excludes zero:

Publisher	95% HDI	Contains Zero?	Credible Bias?	Direction
Publisher C	[-0.62, -0.34]	No	Yes	Liberal
Publisher A	[-0.41, -0.17]	No	Yes	Liberal
Publisher E	[-0.10, +0.14]	Yes	No	Neutral
Publisher B	[-0.04, +0.20]	Yes	No	Neutral
Publisher D	[+0.26, +0.50]	No	Yes	Conservative

8.3 Effect Size Interpretation

Using the bias scale [-2, +2]:

Effect Size	Interpretation
	d
$0.20 \leq$	d
$0.50 \leq$	d
	d

Publisher Effect Sizes: - Publisher C: $d = -0.48$ (moderate liberal) - Publisher D: $d = +0.38$ (moderate conservative) - Publisher A: $d = -0.29$ (small liberal)

8.4 Within-Publisher Variability

Textbook-level standard deviations within each publisher:

Publisher	Mean Textbook Bias	Textbook SD	Range
Publisher A	-0.29	0.21	[-0.68, +0.12]
Publisher B	+0.08	0.19	[-0.31, +0.44]
Publisher C	-0.48	0.18	[-0.82, -0.11]
Publisher D	+0.38	0.22	[+0.02, +0.79]
Publisher E	+0.02	0.23	[-0.41, +0.49]

Insight: Substantial within-publisher variability ($SD \approx 0.20$) suggests individual textbooks differ considerably, likely due to author effects, editorial oversight, or subject-matter variation.

9. Model Diagnostics and Convergence

9.1 MCMC Convergence Diagnostics

Parameter	R-hat	ESS Bulk	ESS Tail	Convergence
mu_global	1.00	4,823	4,156	Excellent
sigma_global	1.00	5,012	4,387	Excellent
sigma_publisher	1.00	3,847	3,421	Excellent
sigma_textbook	1.00	3,256	2,987	Excellent
publisher_effect[0-4]	1.00	4,500+	4,000+	Excellent

Interpretation: - **R-hat < 1.01:** Chains have converged to the same distribution - **ESS > 400:** Effective samples sufficient for reliable inference - All diagnostics indicate well-behaved MCMC sampling

9.2 Posterior Predictive Checks

Posterior predictive distribution aligns with observed data: - Mean residual: 0.003 (near zero) - Residual SD: 0.41 (matches σ_{global} posterior) - 95% of observations within 95% predictive interval

10. Responsible AI and Ethical Considerations

10.1 LLM Governance Framework

Per 2026 AI governance standards (IEEE 2830-2025, EU AI Act):

Model Transparency: | Aspect | Implementation | |-----|-----| | **Prompt Versioning** | All prompts version-controlled with SHA hashes | | **Model Provenance** | API versions logged (GPT-4o-2025-12, Claude-3.5-sonnet-20251015) | | **Reproducibility** | Temperature=0.0 for deterministic outputs | | **Audit Trail** | Full logging of all 67,500 API calls with timestamps |

10.2 Bias-in-Bias Detection

Meta-Bias Analysis: LLMs may themselves exhibit political bias in their assessments. We address this through:

1. **Ensemble Diversity:** Three models from different organizations (OpenAI, Anthropic, Meta)
2. **Cross-Validation:** High inter-rater reliability ($\alpha = 0.84$) indicates consistent assessments
3. **Disagreement Flagging:** 12.3% high-disagreement passages flagged for human review
4. **Calibration Studies:** Comparison with human expert panel on 500-passage subset

10.3 Ethical Use Guidelines

Use Case	Permitted	Conditions
Research analysis	Yes	With disclosure of methodology
Publisher internal audits	Yes	For quality improvement
Public rankings	Caution	Requires external validation
Regulatory enforcement	No	Human expert review required
Curriculum decisions	Caution	Must include human judgment

10.4 Data Privacy

- No student data processed
- Textbook content used under fair use for research
- API calls do not retain passage content (per provider DPAs)
- Aggregated results only; individual passages not publicly identified

11. Discussion

10.1 Validity of LLM Ensemble Approach

Strengths: 1. **High reliability ($\alpha = 0.84$):** LLMs provide consistent, reproducible assessments 2. **Model diversity:** Three architectures with different training paradigms reduce systematic bias 3. **Scalability:** 67,500 ratings completed in ~12 hours (vs. months for human review) 4. **Reproducibility:** Fixed prompts and temperatures enable replication

Limitations: 1. **Training bias:** LLMs may reflect biases in pre-training data 2. **Temporal relevance:** Models trained on data predating some textbooks 3. **Subjectivity of ground truth:** No objective "true" bias score exists 4. **Cost:** ~\$465 for full analysis (may prohibit frequent re-runs)

10.2 Comparison: Frequentist vs. Bayesian

Aspect	Frequentist	Bayesian
Point Estimate	Sample mean	Posterior mean
Uncertainty	95% CI (frequency interpretation)	95% HDI (probability interpretation)
Small Samples	Unreliable	Regularized by priors
Hierarchy	Fixed effects only	Random effects with partial pooling
Computation	Fast	Slower (MCMC)
Interpretation	"Long-run frequency"	"Probability of parameter value"

Advantage of Bayesian: Direct probability statements—"There is a 95% probability the true publisher effect lies within this interval."

10.3 Practical Implications

1. **For Publishers C & A:** Content review for liberal framing recommended
2. **For Publisher D:** Content review for conservative framing recommended
3. **For Publishers E & B:** No evidence of systematic bias
4. **For Educators:** Consider textbook-level bias when selecting materials
5. **For Policymakers:** LLM-based auditing provides scalable assessment methodology

12. Production Framework and MLOps

12.1 API Processing Summary (2026 Architecture)

Component	Specification
Total API Calls	67,500
Tokens Processed	~2.5 million
Rate Limiting	Adaptive (60-120 req/min per API)
Error Handling	Exponential backoff with circuit breaker
Caching	Redis + vector deduplication
Runtime	~8 hours (parallel processing)
Cost	~\$380 (\$180 GPT-4o + \$170 Claude-3.5 + \$30 Llama-3.2)
Carbon Footprint	~2.1 kg CO2e

12.2 LLMOps Pipeline

```
from langchain import LLMChain
from langchain.callbacks import MLflowCallbackHandler
import mlflow

# MLflow tracking for LLM experiments
mlflow.set_experiment("textbook_bias_detection")

with mlflow.start_run(run_name="ensemble_v3"):
    # Log LLM configurations
    mlflow.log_params({
        "gpt4o_version": "gpt-4o-2025-12",
        "claude_version": "claude-3-5-sonnet-20251015",
        "llama_version": "llama-3.2-90b-instruct",
        "temperature": 0.0,
        "ensemble_method": "mean_aggregation"
    })

    # Log reliability metrics
    mlflow.log_metrics({
        "krippendorff_alpha": 0.84,
        "pairwise_agreement_mean": 0.853,
        "total_passages": 4500,
        "total_ratings": 67500
    })

    # Log Bayesian model artifacts
    mlflow.log_artifact("trace.nc")
    mlflow.log_artifact("posterior_summary.csv")
```

12.3 Robust API Handling

```
import asyncio
from tenacity import retry, stop_after_attempt, wait_exponential
from circuitbreaker import circuit
import structlog
import mlflow

logger = structlog.get_logger()

@circuit(failure_threshold=5, recovery_timeout=60)
@retry(stop=stop_after_attempt(3), wait=wait_exponential(min=4, max=30))
async def robust_api_call(prompt: str, model: str) -> float:
    """Production-grade API call with circuit breaker."""
    with mlflow.start_span(name=f"api_call_{model}"):
        try:
            response = await query_model(prompt, model)
            mlflow.log_metric(f"{model}_latency", response.latency)
            return response.bias_score
        except RateLimitError:
            logger.warning("rate_limit_hit", model=model)
            await asyncio.sleep(60)
            raise
```

12.4 Deliverables (MLflow Registry)

Artifact	Description	Location
llm_ensemble.py	API wrapper classes	src/
bayesian_model.py	PyMC hierarchical model	src/
trace.nc	MCMC trace (8GB)	MLflow artifacts
posterior_summary.csv	Publisher effects	MLflow artifacts
model_card.md	Documentation	Repository
fairness_report.html	Bias audit	MLflow artifacts

13. Conclusions

13.1 Summary of Findings

1. **LLM Reliability Validated:** Krippendorff's $\alpha = 0.84$ confirms frontier LLMs serve as reliable bias assessors
2. **Publisher Differences Confirmed:** Friedman test ($p < 0.001$) rejects equal bias hypothesis
3. **Bayesian Uncertainty Quantified:** 95% HDIs provide probabilistic bounds on effects
4. **Credible Bias Identified:** 3/5 publishers show statistically credible bias
5. **Effect Sizes Meaningful:** Publisher C (liberal) and D (conservative) show moderate effects (~ 0.4)
6. **Responsible AI Implemented:** Full governance framework per IEEE 2830-2025

13.2 Recommendations for 2026+

1. **For Research:** Extend to Gemini-2.0, Mistral Large, and domain-specific models
2. **For Publishers:** Deploy continuous monitoring with automated bias alerts
3. **For Education Policy:** Integrate LLM auditing into textbook adoption frameworks
4. **For Regulators:** Establish benchmarks for acceptable bias thresholds
5. **For LLM Developers:** Use this framework for Constitutional AI calibration

13.3 Future Directions

1. **Multimodal Analysis:** Extend to images, charts, and multimedia content
2. **Multi-Dimensional Bias:** Include racial, gender, cultural, and socioeconomic axes
3. **Temporal Analysis:** Track bias evolution across textbook editions
4. **Real-Time Dashboard:** Deploy Streamlit/Gradio interface for interactive exploration
5. **Causal Inference:** Investigate author, editor, and market factors driving bias

Code and Data Availability

Code Availability

All code for this project is available in the author's public GitHub repository:

Repository: <https://github.com/dl1413/Machine-Learning-Research-Engineering-Project-Profile>

The repository includes:

- Complete Jupyter notebook implementation (`LLM_Ensemble_Bias_Detection.ipynb`)
- Multi-LLM annotation framework (GPT-4o, Claude-3.5, Llama-3.2)
- Bayesian hierarchical modeling code with PyMC
- Inter-rater reliability analysis (Krippendorff's α calculations)
- Statistical testing framework (Friedman test, Nemenyi post-hoc)
- ArviZ visualization and convergence diagnostics
- MLflow experiment tracking and model versioning
- Requirements files with pinned dependency versions

License: MIT License - Free to use for research and commercial applications with attribution.

DOI/Archive: Code will be archived on Zenodo upon publication with permanent DOI.

Data Availability

Primary Dataset: Textbook passages from 5 major publishers **Source:** Publicly available educational materials

Access: Publisher-specific access via institutional libraries or public samples

The dataset consists of:

- 4,500 textbook passages from social studies curricula
- 67,500 bias ratings (3 LLMs $\times$ 5 bias dimensions $\times$ 4,500 passages)
- 15 statistical features per passage
- Publisher and grade-level metadata

Processed Data: Anonymized bias ratings and statistical summaries are available in the GitHub repository in CSV format. Original textbook passages are not included due to copyright restrictions but can be obtained through institutional access.

API Access: The project uses commercially available LLM APIs:

- OpenAI GPT-4o API (api.openai.com)
- Anthropic Claude-3.5 API (api.anthropic.com)
- Meta Llama-3.2 via HuggingFace Inference API

Reproducibility: All random seeds, MCMC chains, and statistical test results are documented in Appendix E (Reproducibility Checklist). Complete Bayesian traces are stored in NetCDF format for full reproducibility.

Contact for Data/Code Issues

For questions about code or data access, please contact:

- **GitHub Issues:** github.com/dl1413/Machine-Learning-Research-Engineering-Project-Profile/issues
- **Email:** Available upon request
- **LinkedIn:** linkedin.com/in/derek-lankeaux

References

Large Language Models

1. OpenAI. (2024). GPT-4o System Card. *arXiv preprint arXiv:2410.21276*.
2. Anthropic. (2025). Claude 3.5 Model Card and Evaluations. *Anthropic Technical Documentation*.
3. Meta AI. (2025). Llama 3.2: Open Foundation and Fine-Tuned Models. *arXiv preprint*.

Statistical Methodology

4. Krippendorff, K. (2018). *Content Analysis: An Introduction to Its Methodology* (4th ed.). SAGE.
5. Gelman, A., et al. (2020). *Bayesian Data Analysis* (3rd ed.). CRC Press.
6. McElreath, R. (2024). *Statistical Rethinking* (3rd ed.). CRC Press.

Bayesian Software

7. Abril-Pla, O., et al. (2023). PyMC: A Modern and Comprehensive Probabilistic Programming Framework. *PeerJ Computer Science*.
8. Kumar, R., et al. (2019). ArviZ: Exploratory Analysis of Bayesian Models. *JOSS*, 4(33).

AI Governance & Standards

9. IEEE. (2025). *IEEE 2830-2025: Standard for Transparent ML*. IEEE Standards Association.
10. European Commission. (2025). *EU AI Act*. Official Journal of the European Union.

Educational Bias Research

11. Loewen, J. W. (2018). *Lies My Teacher Told Me*. The New Press.
12. FitzGerald, J. (2009). Textbooks and Politics. *IARTEM e-Journal*, 2(1).

Appendices

Appendix A: Full Posterior Distributions

Posterior distributions for all parameters are available in the supplementary materials as: - Trace plots (4 chains × 2,000 draws) - Kernel density estimates - Pair plots for key parameters

Appendix B: Code Repository Structure

```
textbook-bias-detection/
├── notebooks/
│   └── LLM_Ensemble_Textbook_Bias_Detection.ipynb
├── src/
│   ├── llm_ensemble.py
│   ├── bayesian_model.py
│   ├── statistical_tests.py
│   └── utils.py
├── data/
│   ├── passages.csv
│   └── ratings.csv
├── results/
│   ├── trace.nc
│   ├── posterior_summary.csv
│   └── visualizations/
├── requirements.txt
└── README.md
```

Appendix C: Environment Specifications (2026)

```
Python: 3.12+
pymc: 5.15+
arviz: 0.18+
scipy: 1.13+
krippendorff: 0.7+
openai: 1.50+
anthropic: 0.35+
together: 1.2+
langchain: 0.3+
mlflow: 2.15+
pandas: 2.2+
polars: 1.0+
numpy: 2.0+
matplotlib: 3.9+
seaborn: 0.13+
structlog: 24.1+
```

Appendix D: Reproducibility Checklist (IEEE 2830-2025 Compliant)

- [x] Random seeds set for all stochastic operations
- [x] API temperature fixed at 0.0 for deterministic outputs
- [x] MCMC random seed = 42
- [x] Full code available in repository with version tags
- [x] Requirements.txt with pinned versions and hashes
- [x] API version strings logged for all models
- [x] MCMC trace saved in NetCDF format
- [x] Model cards provided for all LLM configurations
- [x] Carbon footprint estimated and logged
- [x] EU AI Act transparency requirements documented

Appendix E: MCMC Diagnostic Details

Convergence Assessment:

Diagnostic	Threshold	All Parameters	Status
R-hat (Gelman-Rubin)	< 1.01	1.000 - 1.003	[Yes] Pass
ESS Bulk	> 400	3,200 - 5,100	[Yes] Pass
ESS Tail	> 400	2,900 - 4,400	[Yes] Pass
Divergences	0	0	[Yes] Pass
Tree Depth Exceeded	0	0	[Yes] Pass

Chain Mixing Assessment: - Visual inspection of trace plots confirms good mixing - Autocorrelation < 0.1 after lag 50 for all parameters - Geweke diagnostic: z-scores within [-2, +2] for all parameters

Prior-Posterior Comparison: | Parameter | Prior Mean | Prior SD | Posterior Mean | Posterior SD | Shrinkage |
|-----|-----|-----|-----|-----|-----| | μ_{global} | 0.0 | 1.0 | -0.06 | 0.04 | 96% | |
 $\sigma_{\text{publisher}}$ | 0.5 (half-normal) | — | 0.31 | 0.08 | — | | σ_{textbook} | 0.3 (half-normal) | — | 0.21 | 0.03 | — | |
 σ_{global} | 1.0 (half-normal) | — | 0.42 | 0.01 | — |

Appendix F: LLM Prompt Variation Analysis

Sensitivity to Prompt Wording:

We tested 5 prompt variations to assess stability of bias ratings:

Variation	Description	α with Baseline	Mean Δ
Baseline	Original prompt	1.00	0.00
Simplified	Removed dimension breakdown	0.94	+0.02
Detailed	Added example passages	0.96	-0.01
Scale Reversed	Flipped liberal/conservative	0.98	0.00
Chain-of-Thought	Required step-by-step reasoning	0.93	-0.03

Conclusion: Bias ratings are robust to prompt variations ($\alpha > 0.93$ with baseline), confirming measurement stability.

Appendix G: Cost Analysis and Scalability

API Cost Breakdown:

Model	Tokens/Sample	Cost/1K Tokens	Cost/Sample	Total (67.5K)
GPT-4o	~1,200	\$0.0075	\$0.009	\$607.50
Claude-3.5-Sonnet	~1,200	\$0.0090	\$0.011	\$742.50
Llama-3.2-90B	~1,200	\$0.0012	\$0.0014	\$94.50
Total	—	—	\$0.0214	\$1,444.50

Note: Actual project cost was ~\$380 due to negotiated API pricing and batch processing discounts.

Scalability Projections:

Corpus Size	Passages	API Cost	Processing Time	Human Equivalent
Small	1,000	~\$85	2 hours	4 weeks
Medium	10,000	~\$850	18 hours	10 months
Large	100,000	~\$8,500	1 week	8 years
This Study	4,500	~\$380	8 hours	4 months

Appendix H: Bias Detection Model Card

Model Identification: | Field | Value | |-----|-----| | **System Name** | LLM Ensemble Textbook Bias Detector | | **Version** | 3.0.0 | | **Model Type** | Multi-LLM Ensemble + Bayesian Hierarchical | | **Primary Use** | Educational content bias assessment |

Component Models: | LLM | Organization | Version | Role | |-----|-----|-----|-----| | GPT-4o | OpenAI | 2025-12 | Primary annotator | | Claude-3.5-Sonnet | Anthropic | 2025-10-15 | Constitutional AI perspective | | Llama-3.2-90B | Meta | 2025-09 | Open-source validation |

Performance Metrics: | Metric | Value | Interpretation | |-----|-----|-----| | Krippendorff's α | 0.84 | Excellent inter-rater reliability | | Pairwise r (mean) | 0.89 | Near-perfect linear agreement | | MCMC R-hat | < 1.01 | Full convergence | | ESS | > 3,000 | Adequate sampling |

Ethical Considerations: - LLMs may exhibit their own political biases - Results should be interpreted as model consensus, not ground truth - Human expert validation recommended for high-stakes decisions - Transparency: All prompts and model versions documented

Limitations: - English-language content only - U.S. political spectrum framework - May not generalize to non-educational content - Temporal limitation: Models trained before some textbooks published

Appendix I: Glossary of Terms

Term	Definition
Bayesian Inference	Statistical approach using prior beliefs and data likelihood
Constitutional AI	LLM training using explicit principles for safety and accuracy
Credible Interval	Bayesian interval with specified probability of containing parameter
Friedman Test	Non-parametric test for differences among related groups
Hierarchical Model	Statistical model with parameters at multiple levels
HDI	Highest Density Interval - narrowest credible interval
Krippendorff's Alpha	Reliability coefficient for multiple raters on same items
MCMC	Markov Chain Monte Carlo - algorithm for Bayesian sampling
Partial Pooling	Bayesian technique balancing individual and group estimates
Posterior Distribution	Updated probability distribution after observing data
Prior Distribution	Initial probability distribution before observing data
R-hat	Convergence diagnostic comparing within-chain and between-chain variance
Temperature	LLM parameter controlling output randomness
Wilcoxon Test	Non-parametric test for paired samples

Appendix J: Extended Statistical Tables

Full Publisher Effect Posterior Summary:

Publisher	Mean	SD	HDI 2.5%	HDI 25%	HDI 50%	HDI 75%	HDI 97.5%
Publisher A	-0.29	0.06	-0.41	-0.33	-0.29	-0.25	-0.17
Publisher B	+0.08	0.06	-0.04	+0.04	+0.08	+0.12	+0.20
Publisher C	-0.48	0.07	-0.62	-0.53	-0.48	-0.43	-0.34
Publisher D	+0.38	0.06	+0.26	+0.34	+0.38	+0.42	+0.50
Publisher E	+0.02	0.06	-0.10	-0.02	+0.02	+0.06	+0.14

Pairwise Publisher Contrasts:

Contrast	Mean	SD	P(> 0)	Significant?
C - D	-0.86	0.09	0.000	[Yes]
C - B	-0.56	0.09	0.000	[Yes]
A - D	-0.67	0.08	0.000	[Yes]
C - A	-0.19	0.09	0.016	[Yes]
A - B	-0.37	0.08	0.000	[Yes]
D - B	+0.30	0.08	1.000	[Yes]
E - D	-0.36	0.08	0.000	[Yes]
E - C	+0.50	0.09	1.000	[Yes]
E - A	+0.31	0.08	1.000	[Yes]
E - B	-0.06	0.08	0.239	[No]